

B561 Assignment 8 Solutions

1 Object Relational Programming

1. **For this problem you can not use arrays. Solution: see .sql file**

Consider the relational schema `Tree(parent int, child int)` representing the schema for storing a rooted tree.¹ A pair of nodes (m, n) is in `Tree` if m is the parent of n in `Tree`. Notice that a node m can be the parent of multiple children but a node n can have at most one parent node.

It should be clear that for each pair of different nodes m and n in `Tree`, there is a unique shortest path of nodes (n_1, \dots, n_k) in `Tree` from m to n provided we interpret the edges in `Tree` as undirected. A good way to think about this path from a node m to a node n is to first consider the *lowest common ancestor node* of m of n in `Tree`. Then the unique path from m to n is the path that is comprised of the path up the tree from m to this common ancestor and then, from this common ancestor, the path down the tree to the node n . (Note that in this path $n_1 = m$ and $n_k = n$.)

Define the *distance* from m to n to be $k - 1$ if (n_1, \dots, n_k) is the unique shortest path from m to n in `Tree`.

Write a PostgreSQL function `distance(m,n)` that computes the distance in `Tree` for any possible pair of different nodes m and n in `Tree`.

For example, if m is the parent of n in `Tree` then `distance(m,n) = 1` because the shortest path from m to n is (m, n) which has length 1. If m is the grandparent of n in `Tree` then `distance(m,n) = 2` since (m, p, n) is the path from m to n where p is the parent of m and p is a child of n . And if m and n have a common grandparent k then `distance(m,n) = distance(m,k) + distance(k,n) = 4`, etc.

2. **For this problem you can use arrays. Solution: see .sql file**

Consider the relation schema `Graph(source int, target int)` representing the schema for storing a directed graph G of edges.

Now let G be a directed graph that is **acyclic**, a graph without cycles.²

A *topological sort* of an acyclic graph is a list of **all** of its nodes (n_1, n_2, \dots, n_k) such that for each edge (m, n) in G , node m occurs before node n in this list.

Write a PostgreSQL program `topologicalSort()` that returns a topological sort of G .

¹We assume that a tree is a connected graph with a finite number of nodes.

²A cycle is a path (v_0, \dots, v_k) where $v_0 = v_k$.

3. For this problem, you can not use arrays. Solution: see .sql file

Consider the following relational schemas. (You can assume that the domain of each of the attributes in these relations is `int`.)

`partSubpart(pid,sid,quantity)`
`basicPart(pid,weight)`

A tuple (p, s, q) is in `partSubPart` if part s occurs q times as a **direct** subpart of part p . For example, think of a car c that has 4 wheels w and 1 radio r . Then $(c, w, 4)$ and $(c, r, 1)$ would be in `partSubPart`. Furthermore, then think of a wheel w that has 5 bolts b . Then $(w, b, 5)$ would be in `partSubpart`.

A tuple (p, w) is in `basicPart` if basic part p has weight w . A basic part is defined as a part that does not have subparts. In other words, the pid of a basic part does not occur in the pid column of `partSubPart`.

(In the above example, a bolt and a radio would be basic parts, but car and wheel would not be basic parts.)

We define the *aggregated weight* of a part inductively as follows:

- (a) If p is a basic part then its aggregated weight is its weight as given in the `basicPart` relation
- (b) If p is not a basic part, then its aggregated weight is the sum of the aggregated weights of its subparts, each multiplied by the quantity with which these subparts occur in the `partSubpart` relation.

Example tables: The following example is based on a desk lamp with pid 1. Suppose a desk lamp consists of 4 bulbs (with pid 2) and a frame (with pid 3), and a frame consists of a post (with pid 4) and 2 switches (with pid 5). Furthermore, we will assume that the weight of a bulb is 5, that of a post is 50, and that of a switch is 3.

Then the `partSubpart` and `basicPart` relation would be as follows:

partSubPart			basicPart	
pid	sid	quantity	pid	weight
1	2	4	2	5
1	3	1	4	50
3	4	1	5	3
3	5	2		

Then the aggregated weight of a lamp is $4 \times 5 + 1 \times (1 \times 50 + 2 \times 3) = 76$.

Write a PostgreSQL function `aggregatedWeight(p integer)` that returns the aggregated weight of a part p .

4. **For this problem you need to use arrays. Solution: see .sql file**

Consider the relation schema `document(doc int, words text[])` representing a relation of pairs (d, W) where d is a unique id denoting a document and W denotes the set of words that occur in d .

Let \mathbf{W} denote the set of all words that occur in the documents and let t be a positive integer denoting a *threshold*.

Let $X \subseteq \mathbf{W}$. We say that X is t -frequent if

$$\text{count}(\{d \mid (d, W) \in \text{document and } X \subseteq W\}) \geq t$$

In other words, X is t -frequent if there are at least t documents that contain all the words in X .

Write a PostgreSQL program `frequentSets(t int)` that returns the set of all t -frequent set.

In a good solution for this problem, you should use the following rule: if X is not t -frequent then any set Y such that $X \subseteq Y$ is not t -frequent either. In the literature, this is called the *Apriori* rule of the frequent itemset mining problem. This rule can be used as a pruning rule. In other words, if you have determined that a set X is not t -frequent then you no longer have to consider any of X 's supersets.

To learn more about this problem you can visit the site
https://en.wikipedia.org/wiki/Apriori_algorithm.

5. **For this problem you can use arrays. Solution: see .sql file**

For this problem, first read about the k -means clustering problem in

<http://stanford.edu/~cpiech/cs221/handouts/kmeans.html>

Look at the k -means clustering algorithm described in this document. Your task is to implement this algorithm in PostgreSQL for a dataset that consists of a set of points in a 2-dimensional space.

Assume that the dataset is stored in a ternary relation with schema

`Points(p int, x float, y float)`

where p is an integer uniquely identifying a point (x, y) .

Write a PostgreSQL program `kMeans(k integer)` that returns a set of k points that denote the centroids for the points in `Points`. Note that k is an input parameter to the `kMeans` function.

You will need to reason about how to determine when the algorithm terminates. A possible termination condition is to set a number of iterations that denotes how many iterations are run to approximate the centroids. Another termination condition is to consider when the set of centroids no longer changes.

2 Physical Database Organization and Algorithms

6. Consider the following parameters:

block size	=	4096 bytes
block-address size	=	9 bytes
block access time	=	10 ms (micro seconds)
record size	=	200 bytes
record key size	=	12 bytes

Assume that there is a B⁺-tree, adhering to these parameters, that indexes 10⁸ million records on their primary key values.

Show all the intermediate computations leading to your answers.

- (a) Specify (in ms) the minimum time to retrieve a record with key k in the B⁺-tree provided that there is a record with this key.

Solution: We first need to determine the order n of the B⁺-tree. This can be done by finding the largest n such that

$$9(n + 1) + 12n \leq 4096$$

I.e., $n \leq \frac{4087}{21}$. Thus $n = 194$.

The minimum time will be $(\lceil \log_{195}(10^8) \rceil + 1) * 10ms = 50ms$. The +1 occurs because we need one more block access to the record in the data file.

Observe that the minimum time is determined by the largest possible fanout of the nodes in the B⁺-tree. This fanout is maximally $n + 1 = 194 + 1 = 195$.

- (b) Specify (in ms) the maximum time to retrieve a record with key k in the B⁺-tree.

Solution: The maximum time results when we have the minimum branching factor, i.e. $\lceil \frac{194}{2} \rceil + 1 = 98$.


Since the minimum branching factor at the root is 2, we must determine the height of a tree that has half of the nodes 2, i.e., $\frac{10^8}{2} = 5 \times 10^7$ nodes.

The height of such a the tree will be $\lceil \log_{98}(5 \times 10^7) \rceil = 4$.

We then also need one more block access to get the record, so the maximum time is $(4 + 1 + 1)10ms = 60ms$.


So we note that the minimum and maximum times are off by just 1 block access.

7. **Solution: see .pdf file** Consider the following B⁺-tree of order 2 that holds records with keys 2, 7, 9, and 11. (Observe that (a) an internal node of a B⁺-tree of order 2 can have either 1 or 2 keys values, and 2 or 3 sub-trees, and (b) a leaf node can have either 1 or 2 key values.)



- (a) Show the contents of your B⁺-tree after inserting records with keys 6, 10, 14, and 4, in that order.
- (b) Starting from your answer in question 7a, show the contents of your B⁺-tree after deleting records with keys 2, 14, 4, and 10, in that order.

8. **Solution: see .pdf file** Consider an extensible hashing data structure wherein (1) the initial global depth is set at 1 and (2) all directory pointers point to the same **empty** block which has local depth 0. So the hashing structure looks like this:



Assume that a block can hold at most two records.

- (a) Show the state of the hash data structure after each of the following insert sequences:³
- i. records with keys 2 and 6.
 - ii. records with keys 1 and 7.
 - iii. records with keys 4 and 8.
 - iv. records with keys 0 and 9.
- (b) Starting from the answer you obtained for Question 8a, show the state of the hash data structure after each of the following delete sequences:
- i. records with keys 1 and 2.
 - ii. records with keys 6 and 7.
 - iii. records with keys 0 and 9.
9. Let $R(A, B)$ and $S(B, C)$ be two relations and consider their natural join $R \bowtie S$.

Assume that R has 1,500,000 records and that S has 5,000 records.

Furthermore, assume that 30 records of R can fit in a block and that 10 records of S can fit in a block.

Assume that you have a main-memory buffer with 101 blocks.

- (a) How many block IO's are necessary to perform $R \bowtie S$ using the block nested-loops join algorithm? Show your analysis.

Solution: The complexity is $B(S) + \frac{B(R)B(S)}{100}$.

$$B(R) = \frac{1500000}{30} = 50000$$

$$B(S) = \frac{5000}{10} = 500$$

$$\text{So } B(S) + \frac{B(R)B(S)}{100} = 500 + 5.50000 = 250500.$$

³You should interpret the key values as bit strings of length 4. So for example, key value 7 is represented as the bit string 0111 and key value 2 is represented as the bit string 0010.

- (b) How many block IO's are necessary to perform $R \bowtie S$ using the sort-merge join algorithm? Show your analysis.

Solution In this case, we first need to determine how much time is involved in sorting R and S .

External sorting R takes $2B(R) \log_{100}(B(R))$. This is $2 \cdot 50000 \lceil \log_{100}(5 \cdot 10^4) \rceil = 300000$.

Sorting S takes $2B(S) \log_{100}(B(S))$. This is $2 \cdot 500 \lceil \log_{100}(500) \rceil = 2000$.

The merge phase of these sorted files is $B(R) + B(S)$ which is $50000 + 500 = 50500$.

So the total is $300000 + 2000 + 50500 = 352500$.

- (c) Repeat question 9b under the following assumptions.

Assume that there are p different B -values and that these are uniformly distributed in R and S .

Observe that to solve this problem, depending on p , it may be necessary to perform a block nested-loop join per occurrence of a B -value.

Solution:

The time to sort R and S remains the same. So this account for $300000 + 2000 = 302000$.

- Consider the case $p = 1$. Clearly, neither R nor S fit in the buffer. So, we do a block nested-loop join. Since S is the smaller relation and since $p = 1$, this requires $B(S) + \frac{B(R)B(S)}{100}$ block accesses. So, in total, an additional 250500 block accesses.
- Consider the case $p = 2$. Since now we have 2 B -values, we need 2 block nested-loop joins between a file of size $\frac{50000}{2}$ and a file of size $\frac{500}{2}$. Since S is the smaller relation, we can accomplish each of these in time $250 + \frac{250 \cdot 25000}{100} = 62750$. Does together we have $2 \cdot 62750 = 125500$ block accesses.
- Consider the case $p = 3$. Since now we have 3 B -values, we need 3 block nested-loop joins between a file of size $\frac{50000}{3}$ and a file of size $\frac{500}{3}$, so approximately between a file of size 17000 and one of size 170. Since S is the smaller relation, we can accomplish each of these in time $170 + \frac{170 \times 17000}{100} = 29070$. Thus, we need $3 \times 29070 = 87210$ block accesses.
- We leave the case for $p = 4$ as an exercise.
- Consider the case $p = 5$. Since now we have 5 B -values, we need 5 block nested-loop joins between a file of size $\frac{50000}{5} = 10000$ and a file of size $\frac{500}{5} = 100$. Since now each such chunk of S fits in memory, we get, for each such chunk $100 + 10000 = 10100$ block accesses. Since $p = 5$, we get an extra 50500 block accesses.

The cases for $p > 5$ are similar and we leave these as an exercise.

- (d) How many block IO's are necessary to perform $R \bowtie S$ using the hash-join algorithm? Show your analysis.

Solution: Note: For questions 9b and 9d you can assume that the buffer is sufficiently large to hold all records from R and S for any give B -value.

The complexity is $3(B(R) + B(S)) = 3(50000 + 50) = 150150$.

3 Concurrency Control

10. State which of the following schedules S_1 , S_2 , and S_3 over transactions T_1 , T_2 , and T_3 are conflict-serializable, and for each of the schedules that is serializable, given a serial schedule with which that schedule is conflict-equivalent.

(a) $S_1 = R_1(x)R_2(y)R_1(z)R_2(x)R_1(y)$.

This schedule consists of just read operations. Thus, there are no conflicting pairs and thus all T_1 operations can be swapped with all T_2 operations and as such the schedule S_1 can be transformed to the serial schedule $R_1(x)R_1(z)R_1(y)R_2(y)R_2(x)$, i.e. the schedule $T_1; T_2$.

(b) $S_2 = R_1(x)W_2(y)R_1(z)R_3(z)W_2(x)R_1(y)$.

The precedence graph $P(S_2) = \{(T_1, T_2), (T_2, T_1)\}$ which is cyclic. Thus, S_2 is not serializable.

(c) $S_3 = R_1(z)W_2(x)R_2(z)R_2(y)W_1(x)W_3(z)W_1(y)R_3(x)$.

The precedence graph $P(S_3) = \{(T_1, T_3), (T_2, T_1), (T_2, T_3)\}$ which is acyclic. Thus, S_3 is serializable.

Using topological sort on this graph, S_3 is equivalent with the serial schedule $T_2; T_1; T_3$.

11. Give 3 transactions T_1 , T_2 , T_3 and a schedule S on these transactions whose precedence graph (i.e. serialization graph) consists of the edges (T_1, T_2) , (T_2, T_1) , (T_1, T_3) , (T_3, T_2) .

$$\begin{aligned} T_1 &= R_1(u)R_1(x)R_1(z) \\ T_2 &= W_2(u)W_2(x)W_2(w) \\ T_3 &= W_3(w)W_3(z) \end{aligned}$$

$$S = W_2(u)R_1(u)R_1(x)R_1(z)W_2(x)W_3(w)W_3(z)W_2(w).$$

12. Give 3 transactions T_1 , T_2 , and T_3 that each involve read and write operations and a schedule S that is conflict-equivalent with **all** serial schedules over T_1 , T_2 , and T_3 .

We could pick the following transactions

$$\begin{aligned} T_1 &= R_1(x)W_1(x) \\ T_2 &= R_2(y)W_2(y) \\ T_3 &= R_3(z)W_3(z) \end{aligned}$$

In this example there are no conflicting pairs and therefore any schedule over T_1 , T_2 , and T_3 can be transformed into any serial schedule using appropriate swap operations of non-conflicting operations.

13. Consider the following transactions:

```
T1:  read(A);
      read(B);
      if A = 0 then B := B+1;
      write(B).
```

```
T2:  read(B);
      read(A);
      if B = 0 then A := A+1;
      write(A).
```

Let the consistency requirement be $A = 0 \vee B = 0$, and let $A = B = 0$ be the initial values.

- (a) Show that each serial schedule involving transaction T_1 and T_2 preserves the consistency requirement of the database.

Solution:

For the serial schedule $T_1; T_2$ the value for $A = 0$ and that for $B = 1$.

For the serial schedule $T_2; T_1$ the value for $A = 1$ and that for $B = 0$.

- (b) Construct a schedule on T_1 and T_2 that produces a non-serializable schedule.

Solution:

T_1	T_2
R(A)	
	R(B)
	R(A)
	if B = 0 then A := A+1
	W(A)
R(B)	
if A = 0 then B := B+1	
W(B)	

The precedence graph consists of the edges (T_1, T_2) and (T_2, T_1) . This is a cyclic graph so this schedule is not serializable.

- (c) Is there a non-serial schedule on T_1 and T_2 that produces a serializable schedule. If so, give an example.

Solution:

The answer is no.

Assume that the schedule begins with $R_1(A)$. Observe that this action conflicts with $W_2(A)$.

Assume that we consider the partial schedule $R_1(A)R_2(B)$. Observe that $R_2(B)$ conflicts with $W_1(B)$. So if we start the schedule with

$R_1(A)R_2(B)$ we get a cyclic precedence graph. Consequently, if we begin the schedule with $R_1(A)$ then the next action must be $R_1(B)$. So we must have $R_1(A)R_1(B)$. We could now consider $R_2(B)$. But this will again create a cycle in the precedence graph. Thus we are required to do " $R_1(A)R_1(B)$ if $A = 0$ then $B := B + 1$ ". Given this if we consider $R_2(B)$ we again have a problem since we will get a cyclic precedence graph. Thus we conclude that if we begin our schedule with $R_1(A)$, then we must execute the serial schedule $T_1; T_2$.

Now assume that we begin the schedule with $R_2(B)$. A similar analysis as above shows that if we want a serializable schedule, then we need to execute the serial schedule $T_2; T_1$.